

Student Name: Ronaldlee Ejalu

CSC 583

FinalProject

```
!pip install transformers datasets simpletransformers
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Requirement already satisfied: transformers in /usr/local/lib/python3.10/dist-packages (4.31.0.dev0)
Requirement already satisfied: datasets in /usr/local/lib/python3.10/dist-packages (2.12.0)
Requirement already satisfied: simpletransformers in /usr/local/lib/python3.10/dist-packages (0.63.11)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from transformers) (3.12.0)
Requirement already satisfied: huggingface-hub<1.0,>=0.14.1 in /usr/local/lib/python3.10/dist-packages (from transformers) (0.14.1)
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-packages (from transformers) (1.22.4)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from transformers) (23.1)
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.10/dist-packages (from transformers) (6.0)
Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.10/dist-packages (from transformers) (2022.10.31)
Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from transformers) (2.27.1)
Requirement already satisfied: tokenizers!=0.11.3,<0.14,>=0.11.1 in /usr/local/lib/python3.10/dist-packages (from transformers) (0.13.3)
Requirement already satisfied: safetensors>=0.3.1 in /usr/local/lib/python3.10/dist-packages (from transformers) (0.3.1)
Requirement already satisfied: tqdm>=4.27 in /usr/local/lib/python3.10/dist-packages (from transformers) (4.65.0)
Requirement already satisfied: pyarrow>=8.0.0 in /usr/local/lib/python3.10/dist-packages (from datasets) (9.0.0)
Requirement already satisfied: dill<0.3.7,>=0.3.0 in /usr/local/lib/python3.10/dist-packages (from datasets) (0.3.6)
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from datasets) (1.5.3)
Requirement already satisfied: xxhash in /usr/local/lib/python3.10/dist-packages (from datasets) (3.2.0)
Requirement already satisfied: multiprocessing in /usr/local/lib/python3.10/dist-packages (from datasets) (0.70.14)
Requirement already satisfied: fsspec[http]>=2021.11.1 in /usr/local/lib/python3.10/dist-packages (from datasets) (2022.11.0)
Requirement already satisfied: aiohttp in /usr/local/lib/python3.10/dist-packages (from datasets) (3.8.4)
Requirement already satisfied: responses<0.19 in /usr/local/lib/python3.10/dist-packages (from datasets) (0.18.0)
Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-packages (from simpletransformers) (1.10.1)
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.10/dist-packages (from simpletransformers) (1.2.2)
Requirement already satisfied: seqeval in /usr/local/lib/python3.10/dist-packages (from simpletransformers) (1.2.2)
Requirement already satisfied: tensorboard in /usr/local/lib/python3.10/dist-packages (from simpletransformers) (2.11.2)
Requirement already satisfied: wandb>=0.10.32 in /usr/local/lib/python3.10/dist-packages (from simpletransformers) (0.12.10)
Requirement already satisfied: streamlit in /usr/local/lib/python3.10/dist-packages (from simpletransformers) (1.23.1)
Requirement already satisfied: sentencepiece in /usr/local/lib/python3.10/dist-packages (from simpletransformers) (0.1.96)
Requirement already satisfied: attrs>=17.3.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (23.1.0)
Requirement already satisfied: charset-normalizer<4.0,>=2.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (3.1.0)
```

```

Requirement already satisfied: multidict<7.0,>=4.5 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets)
Requirement already satisfied: async-timeout<5.0,>=4.0.0a3 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets)
Requirement already satisfied: yarl<2.0,>=1.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (1.8.2)
Requirement already satisfied: frozenlist>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets)
Requirement already satisfied: aiosignal>=1.1.2 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets)
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.10/dist-packages (from huggingface-hub->datasets)
Requirement already satisfied: urllib3<1.27,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests->transformers)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests->transformers)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests->transformers)
Requirement already satisfied: Click!=8.0.0,>=7.0 in /usr/local/lib/python3.10/dist-packages (from wandb>=0.10.32->sentencepiece)
Requirement already satisfied: GitPython!=3.1.29,>=1.0.0 in /usr/local/lib/python3.10/dist-packages (from wandb>=0.10.32->sentencepiece)
Requirement already satisfied: psutil>=5.0.0 in /usr/local/lib/python3.10/dist-packages (from wandb>=0.10.32->sentencepiece)
Requirement already satisfied: sentry-sdk>=1.0.0 in /usr/local/lib/python3.10/dist-packages (from wandb>=0.10.32->sentencepiece)
Requirement already satisfied: docker-pycreds>=0.4.0 in /usr/local/lib/python3.10/dist-packages (from wandb>=0.10.32->sentencepiece)
Requirement already satisfied: pathtools in /usr/local/lib/python3.10/dist-packages (from wandb>=0.10.32->sentencepiece)
Requirement already satisfied: setproctitle in /usr/local/lib/python3.10/dist-packages (from wandb>=0.10.32->sentencepiece)
Requirement already satisfied: setuptools in /usr/local/lib/python3.10/dist-packages (from wandb>=0.10.32->sentencepiece)
Requirement already satisfied: appdirs>=1.4.3 in /usr/local/lib/python3.10/dist-packages (from wandb>=0.10.32->sentencepiece)
Requirement already satisfied: protobuf!=4.21.0,<5,>=3.19.0 in /usr/local/lib/python3.10/dist-packages (from wandb>=0.10.32->sentencepiece)
Requirement already satisfied: python-dateutil>=2.8.1 in /usr/local/lib/python3.10/dist-packages (from pandas->datasets)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas->datasets) (2022.7.1)
Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from scikit-learn->sentencepiece)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn->sentencepiece)
Requirement already satisfied: altair<6,>=4.0 in /usr/local/lib/python3.10/dist-packages (from streamlit->sentencepiece)
Requirement already satisfied: blinker<2,>=1.0.0 in /usr/local/lib/python3.10/dist-packages (from streamlit->sentencepiece)

```

```

import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from datasets import Dataset

```

```

# mount the drive
from google.colab import drive
drive.mount('gdrive')

```

Drive already mounted at gdrive; to attempt to forcibly remount, call drive.mount("gdrive", force_remount=True).

```

# change the drive to my working directory

```

```
%cd gdrive/MyDrive/CSC583
```

```
/content/gdrive/MyDrive/CSC583
```

▼ Load the data

```
df_raw = pd.read_csv('data/FakeCovid_July2020.csv', encoding='utf-8', sep=',')  
df_raw.head()
```

	ID	ref_category_title	ref_url	pageid	verifiedby	country	class	title
0	FC1	FALSE: The coronavirus is an amplified bacteria	https://www.poynter.org/?ifcn_misinformation=t...	https://www.poynter.org/ifcn-covid-19-misinfor...	La Silla Vacía	Colombia	FALSE	The coronavirus is an amplified bacteria

```
# filter out the entries that are in the English language
```

```
cond = df_raw['lang'] == 'en'
```

```
df_raw = df_raw[cond]
```

```
df_raw.head(3)
```

	ID	ref_category_title	ref_url	pageid	verifiedby	country	class	
2	FC3	False: Chinese converting to Islam after realizing	https://www.poynter.org/?ifcn_misinformation=c...	https://www.poynter.org/ifcn-covid-19-misinfor...	FactCrescendo	India	False	C con' tc
6	FC7	MISLEADING: Captions on a reuploaded video about	https://www.poynter.org/?ifcn_misinformation=c...	https://www.poynter.org/ifcn-covid-19-misinfor...	VERA Files	Philippines	MISLEADING	Capti reupl videc the
8	FC9	Mostly True: Ghana has 307 ambulances with mob...	https://www.poynter.org/?ifcn_misinformation=g...	https://www.poynter.org/ifcn-covid-19-misinfor...	GhanaFact	Ghana	Mostly True	Ghai ambu with ve

```
df_raw.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 2845 entries, 2 to 7621
Data columns (total 19 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ID                     2845 non-null   object
1   ref_category_title     2580 non-null   object
2   ref_url                2845 non-null   object
3   pageid                 2580 non-null   object
4   verifiedby             2845 non-null   object
5   country                2580 non-null   object
6   class                  2845 non-null   object
7   title                  2845 non-null   object
8   published_date         2845 non-null   object
9   country1               2580 non-null   object
10  country2               119 non-null    object
11  country3               24 non-null     object
12  country4               2 non-null      object
13  article_source         2845 non-null   object
14  ref_source             2845 non-null   object
15  source_title           2809 non-null   object
16  content_text           2845 non-null   object
17  category               265 non-null    object
18  lang                   2845 non-null   object
dtypes: object(19)
memory usage: 444.5+ KB

```

```
df_raw['class'] = df_raw['class'].str.lower()
```

```
df_raw['class'].value_counts()
```

```

false                2266
misleading           233
mostly false         117
mixture              41
true                 34
partly false         26
news                 23
mostly true          15
correct attribution   12

```

half true	11
miscaptioned	10
unproven	10
no evidence	9
misattributed	7
(org. doesn't apply rating)	7
labeled satire	5
two pinocchios	3
scam	3
partly true	2
not true	2
partially false	1
fake	1
pants on fire	1
in dispute	1
partially correct	1
unlikely	1
collections	1
mixed	1
partially true	1

Name: class, dtype: int64

```
conditions = [ df_raw['class'] == 'false', df_raw['class'] == 'true', \
               df_raw['class'] == 'mostly true', df_raw['class'] == 'mostly false', \
               df_raw['class'] == 'partly true', df_raw['class'] == 'partly false', \
               df_raw['class'] == 'half true', df_raw['class'] == 'partially false', \
               df_raw['class'] == 'partially true', df_raw['class'] == 'partially correct', \
               df_raw['class'] == 'mixed', \
               df_raw['class'] != 'true', df_raw['class'] != 'false'\
             ]
```

```
choices = ['FALSE','TRUE', 'TRUE', 'FALSE', 'TRUE','FALSE','TRUE' , 'FALSE', 'TRUE', 'TRUE', 'TRUE', 'UNKNOWN', 'UNKNOWN']
df_raw['class'] =np.select(conditions, choices, default=np.nan)
```

```
df_raw['class'].value_counts()
```

FALSE	2410
UNKNOWN	370

```

    TRUE      65
    Name: class, dtype: int64

df_raw.shape

(2845, 19)

# manual one hot encoding of the class values into 0, 1 and 2
cond_label = [df_raw['class'] == 'FALSE', \
               df_raw['class'] == 'TRUE', \
               df_raw['class'] == 'UNKNOWN'
               ]

choices_label = [0, 1, 2]
df_raw['label'] = np.select(cond_label, choices_label, default=np.nan)
# change the type of the label column to an integer
df_raw['label'] = df_raw['label'].astype('int')

# let's extract out the attributes of interest
df_raw_ds = df_raw[['ID', 'label', 'class', 'content_text']]

df_raw_ds['label'].value_counts()

0      2410
2       370
1         65
Name: label, dtype: int64

df_raw_ds.head()

```

	ID	label	class	content_text
2	FC3	0	FALSE	The fact behind every news!, Ever since the Wo...

```

# separate the class from the data frame columns
labels_np = df_raw['class'].values

# apply one hot encoding
labels_one_hot_enc = pd.get_dummies(labels_np)
labels_one_hot_enc.shape

(2845, 3)

# convert the one encoded labels into a numpy array
labels_one_hot_np = labels_one_hot_enc.to_numpy()
print('Shape of the numpy array: %s'%str(labels_one_hot_np.shape))
labels_one_hot_np

Shape of the numpy array: (2845, 3)
array([[1, 0, 0],
       [0, 0, 1],
       [0, 1, 0],
       ...,
       [1, 0, 0],
       [1, 0, 0],
       [1, 0, 0]], dtype=uint8)

# drop the class column since it is now encoded
df_raw.drop('class', axis=1, inplace = True)

# separate the content_text from the data frame
content_texts = df_raw['content_text'].values
type(content_texts)

numpy.ndarray

```



```
# lets extract the labels and text separately
label = df_raw_ds['label']
content_texts = df_raw_ds['content_text']
```

```
label.shape
```

```
(2845,)
```

```
content_texts.shape
```

```
(2845,)
```

▼ Stratified Sampling

```
X_train, X_test, y_train, y_test = train_test_split(content_texts, label, test_size=0.5, random_state=1, stratify=label)
```

```
y_train.value_counts(normalize=True) * 100
```

```
0    84.739803
2    13.009845
1     2.250352
Name: label, dtype: float64
```

```
y_test.value_counts(normalize=True) * 100
```

```
0    84.680253
2    13.000703
1     2.319044
Name: label, dtype: float64
```

Both our target values of 2 (Unknown) and 1 (True) of the train and test sets balance out.

```
# convert the series into a list
```

```
X_train_L = X_train.tolist()
X_test_L = X_test.tolist()
y_train_L = y_train.tolist()
y_test_L = y_test.tolist()
```

▼ Convert the list into a list of dictionary types

```
# convert the training list into a list of dictionary type
X_train_List_of_dict = []
for key, value in zip(y_train_L, X_train_L):
    data = {
        "label" : key,
        "text" : value
    }
    #print(type(key), type(value))
    X_train_List_of_dict.append(data)

# convert the test list into a list of dictionary type
X_test_List_of_dict = []
for key, value in zip(y_test_L, X_test_L):
    data = {
        "label" : key,
        "text" : value
    }
    X_test_List_of_dict.append(data)
```

▼ let's load a data set into a format that suits our model

verification of the list of dictionaries

```
for i in range(len(X_train_List_of_dict) - 1420):
```

```
print(X_train_List_of_dict[i])
```

```
{'label': 0, 'text': ', A MESSAGE IS being shared across Facebook and WhatsApp claiming that criminals are handing out  
{\'label\': 1, \'text\': \'Speak Now, The Covid-19 pandemic has claimed the lives of several health workers across the world
```

```
X_train_set = Dataset.from_list(X_train_List_of_dict)
```

```
X_test_set = Dataset.from_list(X_test_List_of_dict)
```

```
# convert both the training and test set into torch format
```

```
X_train_set = X_train_set.with_format('torch')
```

```
X_test_set = X_test_set.with_format('torch')
```

```
# look at the 100th element from the training data set
```

```
print(X_train_set[100])
```

```
{'label': tensor(0), 'text': '"The post's caption reads: "Good News! Wuhans CORONA VIRUS can be cured By one bowl of fr
```

```
# look at the 100th element from the test data set
```

```
print(X_test_set[100])
```


```
{'label': tensor(0), 'text': 'According to News NT, microbiologist Elisa Granato died two days after she was administer
```

```
from datasets import DatasetDict
```

```
datasets = DatasetDict(  
    {  
        'train' : X_train_set, \  
        'test' : X_test_set  
    }  
)
```

```
print(datasets['train'][100])
```

```
{'label': tensor(0), 'text': '"The post's caption reads: "Good News! Wuhans CORONA VIRUS can be cured By one bowl of fr
```



```
from datasets import ClassLabel
import random
from IPython.display import display, HTML
```

Disclaimer: I borrowed this line of code from Prof. Noriko

```
def show_random_elements(dataset, num_examples=10):
    assert num_examples <= len(dataset), "Can't pick more elements than there are in the dataset."
    picks = []
    for _ in range(num_examples):
        pick = random.randint(0, len(dataset)-1)
        while pick in picks: # to avoid duplicates
            pick = random.randint(0, len(dataset)-1)
        picks.append(pick)

    df = pd.DataFrame(dataset[picks])

    for column, typ in dataset.features.items():
        print("column={}, type={}".format(column, typ))
        if isinstance(typ, ClassLabel):
            df[column] = df[column].transform(lambda i: typ.names[i])
    display(HTML(df.to_html()))

show_random_elements(datasets["train"])
```

```
column=label, type=Value(dtype='int64', id=None)
column=text, type=Value(dtype='string', id=None)
```

label

text

It features one photo of an individual wearing a white gown administering an injection into the arm of another person dressed in military fatigues. It also features a second image of seven people dressed in military fatigues and masks. Below is a screenshot of the misleading post on Facebook: The English and Chinese text in the image reads: "First shot of Covid-19 vaccine developed by China. The postgraduate research team leader took the first shot. The world's first new coronavirus vaccine was injected into the left arm of the lead researcher, Dr Chen Wei, today. Seven medical volunteers, all members of the Communist Party also received the new Coronavirus vaccine. State Council Vice Premier Sun Chunlan visited Chen Wei and the expertise team members.", COVID-19 was first detected in the Chinese city of Wuhan in December 2019. It has since spread worldwide, killing at least 13,000 and infecting 300,000 others, according to this AFP report dated March 22, 2020. The image was also shared here, here, here, here and here on Facebook, here, here, here and here on Twitter and here on Instagram, alongside a similar claim. The claim is misleading. Chinese state media organisation The People's Daily published this tweet on March 5, 2020 which labels the claim "FAKENEWS", alongside a link to this post on its Facebook page. Did Chinese PLA medics test a coronavirus vaccine on themselves? #FAKENEWS. Clinical trials remain underway and this is what the real story is: <https://t.co/lxrdGlerZt> pic.twitter.com/G9BwnEHsvD, The Facebook post reads, in part: "PLA medics are not testing #COVID19 vaccines on themselves, as some would have you believe.", "Photos of a female PLA medic getting a shot have gone viral on the Internet and have made many people wonder if the Chinese PLA is testing a coronavirus vaccine on themselves... the human clinical trial assumption is simply fake news.", "The post was published alongside a photo identical to that seen in the misleading image. The People's Daily states the photo shows Chen Wei, an academician at China's Peoples Liberation Army, receiving a shot for "strengthening immunity" before her departure to Wuhan to join the aid effort.", Below is a screenshot of the Facebook post: Below is screenshot comparison of the photo from the misleading Facebook post (L) and the photo from the People's Daily Facebook post (R): This Global Times report dated March 4, 2020 states Chen was sent to Wuhan, and "achieved a major breakthrough in developing a vaccine to the novel coronavirus". "The report states, in part: "The military medical expert team has worked in Wuhan, the coronavirus epicenter, for more than a month, and the team led by Chen Wei, academician at the Peoples Liberation Army (PLA) Academy of Military Medical Sciences, is seizing every minute and second and focusing all efforts on this emergency scientific research". China announced on March 17, 2020 that it will start human trials for a COVID-19 vaccine manufactured by Chen's team. "This Global Times report published on the same date reads in part: "The coronavirus vaccine developed by Chens team passed the registration review of clinical research and was approved for clinical trial on Monday night.", "After arriving at Wuhan, Chen Wei, a Chinese military major general, and her team launched pharmacy, pharmaceutical and toxicological research on COVID-19 vaccine in cooperation with local companies, based on their previous successful experiences in developing Ebola vaccine.", "Chen, the military scientist, was quoted as saying in the report that the vaccine has been tested on safety, effectiveness and quality control and is prepared for massive production. And the biotechnology firm has started to recruit volunteers for the vaccine.", The World Health Organization (WHO) has said there is currently no medicine that can prevent or treat COVID-19. The Q&A section of its website reads, in part: "To date, there is no vaccine and no specific antiviral medicine to prevent or treat COVID-2019... Possible vaccines and some specific drug treatments are under investigation. They are being tested through clinical trials. WHO is coordinating efforts to develop vaccines and medicines to prevent and treat COVID-19.", Copyright © AFP 2017-2020. All rights reserved. Users can access and consult this website and use the share features available for personal, private, and non-commercial purposes. Any other use, in particular any reproduction, communication to the public or distribution of the content of this website, in whole or in part, for any other purpose and/or by any other means, without a specific licence agreement signed with AFP, is strictly

prohibited. The subject matter depicted or included via links within the Fact Checking content is provided to the extent necessary for correct understanding of the verification of the information concerned. AFP has not obtained any rights from the authors or copyright owners of this third party content and shall incur no liability in this regard. AFP and its logo are registered trademarks.

Read more, Both senators have responded to the coronavirus crisis by doing their jobs as legislators and by providing assistance to frontliners, , , Claim: Senators Grace Poe and Risa Hontiveros have been silent over the coronavirus crisis, but were very vocal about the closure of ABS-CBN., A video posted by YouTube channel Boknoy D Great contained this claim. It was posted on May 14., "Kapag COVID ang pinaguusapan kung saan nakasalalay ang buhay ng mga Pilipino, mga kababayan, ay tila walang pakialam ang mga senador na ito, mga kababayan. Pero 'pag ABS-CBN na ang pinag-uusapan, ay sumusulpot na

▼ create a tokenizer for the whole dataset

halambay na mabuu o mga solusyong dapat gawin maipati sa pamamagitan sa administrasyong Duterte. At inihimisan mga

```
from transformers import AutoTokenizer
tokenizer = AutoTokenizer.from_pretrained('bert-base-cased')

platform, tagged the video for fact checkers to verify. It had 7,049 views as of writing., Rating: 1.0/5.0, The facts. Don't be and

def tokenize_function(examples):
    """Function that tokenizes text"""
    return tokenizer(examples['text'], padding='max_length', truncation=True)

Example 21, which seems to create a center for disease control, while Hontiveros voted against it in 2020, she had Senate Bill

# tokenize the whole data set.
tokenized_datasets = datasets.map(tokenize_function, batched=True)
```

create the train/test datasets and model, and train the model using Trainer

"ripple" (to make waves). Read more. Email: Password: Thank You! In/View your profile page/here. Click close to

```
train_dataset = tokenized_datasets["train"].shuffle(seed=42)
test_dataset = tokenized_datasets["test"].shuffle(seed=42)

print(f'Train shape: {train_dataset.shape}, Test shape: {test_dataset.shape}')

Train shape: (1422, 5), Test shape: (1423, 5)

dogs and things like that, these viruses are transmitted from the animal to the people and that's why China has been the source
```

```
!pip install --upgrade accelerate
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Requirement already satisfied: accelerate in /usr/local/lib/python3.10/dist-packages (0.20.3)
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-packages (from accelerate) (1.22.4)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from accelerate) (23.1)
Requirement already satisfied: psutil in /usr/local/lib/python3.10/dist-packages (from accelerate) (5.9.5)
Requirement already satisfied: pyyaml in /usr/local/lib/python3.10/dist-packages (from accelerate) (6.0)
Requirement already satisfied: torch>=1.6.0 in /usr/local/lib/python3.10/dist-packages (from accelerate) (2.0.1+cu118)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from torch>=1.6.0->accelerate) (3.1)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.10/dist-packages (from torch>=1.6.0->accelerate) (4.5.0)
Requirement already satisfied: sympy in /usr/local/lib/python3.10/dist-packages (from torch>=1.6.0->accelerate) (1.11.1)
Requirement already satisfied: networkx in /usr/local/lib/python3.10/dist-packages (from torch>=1.6.0->accelerate) (3.1)
Requirement already satisfied: Jinja2 in /usr/local/lib/python3.10/dist-packages (from torch>=1.6.0->accelerate) (3.1.2)
Requirement already satisfied: triton==2.0.0 in /usr/local/lib/python3.10/dist-packages (from torch>=1.6.0->accelerate) (2.0.0)
Requirement already satisfied: cmake in /usr/local/lib/python3.10/dist-packages (from triton==2.0.0->torch>=1.6.0->accelerate) (3.25.2)
Requirement already satisfied: lit in /usr/local/lib/python3.10/dist-packages (from triton==2.0.0->torch>=1.6.0->accelerate) (16.0.6)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from Jinja2->torch>=1.6.0->accelerate) (2.1.2)
Requirement already satisfied: mpmath>=0.19 in /usr/local/lib/python3.10/dist-packages (from sympy->torch>=1.6.0->accelerate) (1.3.0)
```

markets permit," he said in an email. "Having said this, we have yet to identify the host animal for the COVID-19 virus. ... It has

We load the model and specify the number of expected labels

cases in China "had some link" to the market, as did numerous news articles highlighted by Brandewie in response to Politifact's

```
from transformers import AutoModelForSequenceClassification
```

```
model = AutoModelForSequenceClassification.from_pretrained("bert-base-cased", num_labels=3)
```

Some weights of the model checkpoint at bert-base-cased were not used when initializing BertForSequenceClassification:

- This IS expected if you are initializing BertForSequenceClassification from the checkpoint of a model trained on another task (e.g. language modeling) with an architecture similar to BertForSequenceClassification.
- This IS NOT expected if you are initializing BertForSequenceClassification from the checkpoint of a model that you expect to be initialized from the checkpoint of a model initialized from the model checkpoint at bert-base-cased and are expecting it to be trained for a sequence classification task (e.g. sentiment classification).

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

possible that the current outbreak was caused by a person in rural China who came into contact with a bat or bat guano and then

we create a TrainingArguments class which contains all the hyperparameters you can tune as well as flags for activating different training options. So, first we specify where to save the checkpoints from our training.

of viruses to be transmitted from animals to humans, but experts said this risk exists across the globe and is not unique to

```
!pip install -q git+https://github.com/huggingface/peft.git git+https://github.com/huggingface/transformers.git
```

```
Installing build dependencies ... done
Getting requirements to build wheel ... done
Preparing metadata (pyproject.toml) ... done
Installing build dependencies ... done
Getting requirements to build wheel ... done
Preparing metadata (pyproject.toml) ... done
```

Our ecological footprint brings us closer to wildlife in remote areas and the wildlife trade brings these animals into urban centers.

```
from transformers import TrainingArguments, Trainer
```

```
training_args = TrainingArguments(output_dir="test_trainer", evaluation_strategy='epoch')
```

Wolves eating the sheep represent the growing farm camps in these forests and images of the hunter hunt translocating wildlife around the world. Wolves raising a lot of domestic livestock that become exposed to viruses through wildlife.

Since Trainer doesn't automatically evaluate model performance during training, we pass Trainer a function to compute metrics. The Evaluate library provides a simple accuracy function you can load with the evaluate.load function.

because the culture where people eat bats and snakes and dogs and things like that, these viruses are transmitted from the

```
!pip install evaluate
```

Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public/simple/>

Collecting evaluate

Downloading evaluate-0.4.0-py3-none-any.whl (81 kB)

81.4/81.4 kB 10.1 MB/s eta 0:00:00

Requirement already satisfied: datasets>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from evaluate) (2.12.0)

Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-packages (from evaluate) (1.22.4)

Requirement already satisfied: dill in /usr/local/lib/python3.10/dist-packages (from evaluate) (0.3.6)

Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from evaluate) (1.5.3)

Requirement already satisfied: requests>=2.19.0 in /usr/local/lib/python3.10/dist-packages (from evaluate) (2.27.1)

Requirement already satisfied: tqdm>=4.62.1 in /usr/local/lib/python3.10/dist-packages (from evaluate) (4.65.0)

Requirement already satisfied: xxhash in /usr/local/lib/python3.10/dist-packages (from evaluate) (3.2.0)

Requirement already satisfied: multiprocessing in /usr/local/lib/python3.10/dist-packages (from evaluate) (0.70.14)

Requirement already satisfied: fsspec[http]>=2021.05.0 in /usr/local/lib/python3.10/dist-packages (from evaluate) (2023.1.0)

Requirement already satisfied: huggingface-hub>=0.7.0 in /usr/local/lib/python3.10/dist-packages (from evaluate) (0.15.1)

Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from evaluate) (23.1)

Requirement already satisfied: responses<0.19 in /usr/local/lib/python3.10/dist-packages (from evaluate) (0.18.0)

Requirement already satisfied: pyarrow>=8.0.0 in /usr/local/lib/python3.10/dist-packages (from datasets>=2.0.0->evaluate) (10.0.1)

Requirement already satisfied: aiohttp in /usr/local/lib/python3.10/dist-packages (from datasets>=2.0.0->evaluate) (3.8.4)

Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.10/dist-packages (from datasets>=2.0.0->evaluate) (6.0.1)

Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from huggingface-hub>=0.7.0->evaluate) (3.12.2)


```
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.10/dist-packages (from huggingface-
Requirement already satisfied: urllib3<1.27,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests>=2.19.0-
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->ev
Requirement already satisfied: charset-normalizer~=2.0.0 in /usr/local/lib/python3.10/dist-packages (from requests>=2.1
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->evaluate
Requirement already satisfied: python-dateutil>=2.8.1 in /usr/local/lib/python3.10/dist-packages (from pandas->evaluate
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas->evaluate) (2022.7.
Requirement already satisfied: attrs>=17.3.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets>=2.0.0-
Requirement already satisfied: multidict<7.0,>=4.5 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets>=
Requirement already satisfied: async-timeout<5.0,>=4.0.0a3 in /usr/local/lib/python3.10/dist-packages (from aiohttp->da
Requirement already satisfied: yarl<2.0,>=1.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets>=2.0.0
Requirement already satisfied: frozenlist>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets>=2.
Requirement already satisfied: aiosignal>=1.1.2 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets>=2.0
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.1->pandas
Installing collected packages: evaluate
Successfully installed evaluate-0.4.0
```

```
import evaluate
```

```
metric = evaluate.load("accuracy")
```

Downloading builder script: 100%

4.20k/4.20k [00:00<00:00, 281kB/s]

We define a function `compute_metrics` that calls `Compute` on `metric` to calculate the accuracy of your predictions. We convert the predictions to logits before passing them to `compute`.

India. Each news story on FACILY is backed by factual evidence/data from official sources that is either available in the public

```
def compute_metrics(eval_pred):
    logits, labels = eval_pred
    predictions = np.argmax(logits, axis=-1)
    return metric.compute(predictions=predictions, references=labels)
```

We create a Trainer object with our model, training arguments, training and test datasets and evaluation function.

“criminally prosecuted” for the deaths of children using the dengue vaccine. **FACTS:** It was not an NIH-made dengue vaccine.

```
trainer = Trainer(model=model, \
```

```
args = training_args, \
train_dataset=train_dataset, \
eval_dataset = test_dataset, \
compute_metrics = compute_metrics
)
```

tested on human subjects in the Philippines to date, based on the records of the U.S. Clinical Trials registry. "Meanwhile

We then fine-tune the model by calling train()

(Metro Manila), where dengue cases were highest that year." The Philippines, Mexico, Brazil, and El Salvador were so far the

trainer.train()

/usr/local/lib/python3.10/dist-packages/transformers/optimization.py:411: FutureWarning: This implementation of AdamW is deprecated. Please use the one in transformers.optimization.AdamW instead.
warnings.warn(

wandb: Logging into wandb.ai. (Learn how to deploy a W&B server locally: <https://wandb.me/wandb-server>)

wandb: You can find your API key in your browser here: <https://wandb.ai/authorize>

wandb: Paste an API key from your profile and hit enter, or press ctrl+c to quit:

wandb: Appending key for api.wandb.ai to your netrc file: /root/.netrc

Tracking run with wandb version 0.15.4

Run data is saved locally in /content/gdrive/MyDrive/CSC583/wandb/run-20230611_030305-zogrlfh3

Syncing run [major-snowball-8](#) to [Weights & Biases](#) (docs)

View project at https://wandb.ai/nlp_ron/huggingface

View run at https://wandb.ai/nlp_ron/huggingface/runs/zogrlfh3

[534/534 09:13, Epoch 3/3]

Epoch	Training Loss	Validation Loss	Accuracy
1	No log	0.479695	0.846803
2	No log	0.495386	0.846803
3	0.510600	0.500322	0.846803

TrainOutput(global_step=534, training_loss=0.5079256586367717, metrics={'train_runtime': 682.9406, 'train_samples_per_second': 6.247, 'train_steps_per_second': 0.782, 'total_flos': 1122441840027648.0, 'train_loss': 0.5079256586367717, 'epoch': 3.0})

2020 Official Gazette President Aquino sneaks at the launching of dengue vaccine school-based immunization in Region III

```
# generate predictions for the test set
predictions = trainer.predict(test_dataset)
```

information on dengue vaccine. Nov. 29, 2020. Doctor, an alleged victim of Dengvaxia. Senate of the Philippines Press Release. [https://www.senate.gov.ph/2020/11/29/doctor-an-alleged-victim-of-dengvaxia/](#)

```

preds = np.array([np.argmax(entry, axis=-1) for entry in predictions[0]])
print(preds)

```

```

[0 0 0 ... 0 0 0]

```

evidence. Find out more about this initiative and our methodology.), TAGS, FOR FURTHER READING, Fact Check Filipino, Jun [https://www.factcheck.ph/2020/06/01/](#)

```

f1_metric = evaluate.load('f1')

```

```

results = f1_metric.compute(predictions=preds, references=predictions[1], average='macro')
print(round(results['f1'], 3))

```

Downloading builder script: 100%

6.77k/6.77k [00:00<00:00, 505kB/s]

0.306

died, and his wife was hospitalized, after the couple ingested the fish tank cleaner to avoid contracting the coronavirus. The [https://www.philippinepostonline.com/2020/06/01/](#)

The macro-average treats all classes equally.

same active ingredient found in anti-malarial drugs that President Donald Trump has hailed as a possible treatment for the novel [https://www.philippinepostonline.com/2020/06/01/](#)

```

results = f1_metric.compute(predictions=preds, references=predictions[1], average='micro')
print(round(results['f1'], 3))

```

0.847

from mainstream news outlets, did not make it clear that the couple ingested fish-tank cleaner, rather than the drug form of [https://www.philippinepostonline.com/2020/06/01/](#)

```

# save the model

```

```

saved_dir = '/content/gdrive/MyDrive/CSC583/data/saved_results'

```

```

tokenizer.save_pretrained(saved_dir)

```

```

model.save_pretrained(saved_dir)

```

for comment by deadline., Officials from Banner Health, the Phoenix hospital system that treated the husband and wife, warned [https://www.philippinepostonline.com/2020/06/01/](#)

```

%ls -l /content/gdrive/MyDrive/CSC583/data/saved_results

```

```

total 424032

```

```

-rw-r--r-- 1 root root      884 Jun 11 03:13 config.json

```

```

-rw-r--r-- 1 root root 433323125 Jun 11 03:13 pytorch_model.bin

```

```

-rw-r--r-- 1 root root      125 Jun 11 03:13 special_tokens_map.json

```

```

-rw-r--r-- 1 root root      315 Jun 11 03:13 tokenizer_config.json

```

```
-rw----- 1 root root    669188 Jun 11 03:13 tokenizer.json
-rw----- 1 root root    213450 Jun 11 03:13 vocab.txt
```

hydroxychloroquine as potential coronavirus treatments, and he tweeted that hydroxychloroquine, combined with an antibiotic

```
from sklearn.metrics import classification_report
print(classification_report(predictions[1], preds))
```

	precision	recall	f1-score	support
0	0.85	1.00	0.92	1205
1	0.00	0.00	0.00	33
2	0.00	0.00	0.00	185
accuracy			0.85	1423
macro avg	0.28	0.33	0.31	1423
weighted avg	0.72	0.85	0.78	1423

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and
_warn_prf(average, modifier, msg_start, len(result))
```

additive may treat coronavirus." March 20, 2020. Banner Health. "Banner Health experts warn against self-medicating to prevent

Since our data set is imbalanced, the weighted average score is 0.78

```
from sklearn.metrics import confusion_matrix
confusion_matrix(predictions[1], preds)
```

```
array([[1205,  0,  0],
       [ 33,  0,  0],
       [ 185,  0,  0]])
```

Write the results to a file

```
outfile = open('true_positives.csv', 'w')
```

```
for target, output, data in zip(predictions[1], preds, test_dataset['text']):
    if target == output:
        outfile.write(f'{output}, {data}')
outfile.close()
```

▼ Data Augmentation

```
!pip install nlpaug
```

Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public/simple/>

Collecting nlpaug

Downloading nlpaug-1.1.11-py3-none-any.whl (410 kB)

410.5/410.5 kB 30.5 MB/s eta 0:00:00

Requirement already satisfied: numpy>=1.16.2 in /usr/local/lib/python3.10/dist-packages (from nlpaug) (1.22.4)
 Requirement already satisfied: pandas>=1.2.0 in /usr/local/lib/python3.10/dist-packages (from nlpaug) (1.5.3)
 Requirement already satisfied: requests>=2.22.0 in /usr/local/lib/python3.10/dist-packages (from nlpaug) (2.27.1)
 Requirement already satisfied: gdown>=4.0.0 in /usr/local/lib/python3.10/dist-packages (from nlpaug) (4.6.6)
 Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from gdown>=4.0.0->nlpaug) (3.12.0)
 Requirement already satisfied: six in /usr/local/lib/python3.10/dist-packages (from gdown>=4.0.0->nlpaug) (1.16.0)
 Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (from gdown>=4.0.0->nlpaug) (4.65.0)
 Requirement already satisfied: beautifulsoup4 in /usr/local/lib/python3.10/dist-packages (from gdown>=4.0.0->nlpaug) (4.11.2)
 Requirement already satisfied: python-dateutil>=2.8.1 in /usr/local/lib/python3.10/dist-packages (from pandas>=1.2.0->nlpaug) (2.8.2)
 Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas>=1.2.0->nlpaug) (2022.7.1)
 Requirement already satisfied: urllib3<1.27,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests>=2.22.0->nlpaug) (1.26.15)
 Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests>=2.22.0->nlpaug) (2022.9.24)
 Requirement already satisfied: charset-normalizer~2.0.0 in /usr/local/lib/python3.10/dist-packages (from requests>=2.22.0->nlpaug) (2.0.12)
 Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests>=2.22.0->nlpaug) (3.4)
 Requirement already satisfied: soupsieve>1.2 in /usr/local/lib/python3.10/dist-packages (from beautifulsoup4->gdown>=4.0.0->nlpaug) (2.3)
 Requirement already satisfied: PySocks!=1.5.7,>=1.5.6 in /usr/local/lib/python3.10/dist-packages (from requests>=2.22.0->nlpaug) (1.7.1)
 Installing collected packages: nlpaug
 Successfully installed nlpaug-1.1.11

```
from tqdm import tqdm
import nlpaug.augmenter.char as nac
import nlpaug.augmenter.word as naw
```

```

import nlpaug.augmenter.sentence as nas
import nlpaug.flow as naf

from nlpaug.util import Action
from sklearn.utils import shuffle

aug = naw.ContextualWordEmbsAug(
    model_path='bert-base-uncased', action="insert")
#augmented_text = aug.augment(text)

```

Downloading (...)okenizer_config.json: 100% 28.0/28.0 [00:00<00:00, 2.06kB/s]

Downloading (...)lve/main/config.json: 100% 570/570 [00:00<00:00, 46.4kB/s]

Downloading (...)solve/main/vocab.txt: 100% 232k/232k [00:00<00:00, 548kB/s]

Downloading (...)main/tokenizer.json: 100% 466k/466k [00:00<00:00, 734kB/s]

Downloading model.safetensors: 100% 440M/440M [00:01<00:00, 417MB/s]

```

def augment_text(df,samples=300,pr=0.2):
    aug.aug_p=pr
    new_text=[]

    ##selecting the minority class samples
    df_n=df[df.label==1].reset_index(drop=True)

    ## data augmentation loop
    for i in tqdm(np.random.randint(0,len(df_n),samples)):

        text = df_n.iloc[i]['text']
        augmented_text = aug.augment(text)
        new_text.append(augmented_text)

    ## dataframe
    new=pd.DataFrame({'text':new_text,'label':1})

```

```
df=shuffle(df.append(new).reset_index(drop=True))

# create a dictionary by passing series as values
dict_f = {'text': X_train, \
          'label': y_train \
          }
# create a train dataframe
df_train = pd.DataFrame(dict_f)
```

```
df_train.shape
```

```
(1422, 2)
```

```
df_train.head(4)
```

	text	label
4537	, A MESSAGE IS being shared across Facebook an...	0
6334	Speak Now, The Covid-19 pandemic has claimed t...	1
318	Responding to comments on his own post, the au...	0
5280	Read more, The quotes in these graphics featur...	0

```
## change samples to 0 for no augmentation
train = augment_text(df_train,samples=1000) # before it was 400
```

```
100%|██████████| 1000/1000 [10:05<00:00, 1.65it/s]
```

```
<ipython-input-68-a526a80b7e7f>:18: FutureWarning: The frame.append method is deprecated and will be removed from pandas
df=shuffle(df.append(new).reset_index(drop=True))
```

```
train['label'].value_counts()
```

```
0    1205
```

```
1    1032
2    185
Name: label, dtype: int64
```

```
y_test.value_counts()
```

```
0    1205
2     185
1      33
Name: label, dtype: int64
```

Separate the labels from text

```
X_train_L = train['text'].tolist()
y_train_L = train['label'].tolist()
```

```
len(X_train_L)
```

```
2422
```

```
len(y_train_L)
```

```
2422
```

▼ Convert the list into a list of dictionary types

```
# convert the training list into a list of dictionary type
X_train_List_of_dict = []
for key, value in zip(y_train_L, X_train_L):
    data = {
        "label" : key,
        "text" : str(value)
```



```
}
X_train_List_of_dict.append(data)
```

▼ let's load the trainind data set into a format that suits our model

```
X_train_set = Dataset.from_list(X_train_List_of_dict)
len(X_train_set)
```

2422

```
# convert the training data set into torch format
X_train_set = X_train_set.with_format('torch')
```

```
# look at the 318th element from the training data set
print(X_train_set[318])
```

```
{'label': tensor(0), 'text': 'The post claims that USAID and Uganda Red Cross Society are undertaking research to ascer
```

```
datasets = DatasetDict(
{
    'train' : X_train_set, \
    'test' : X_test_set
})
```

```
# print the 318th element
print(datasets['train'][318])
```

```
{'label': tensor(0), 'text': 'The post claims that USAID and Uganda Red Cross Society are undertaking research to ascer
```

```
# let's tokenize the whole data set
```

```
tokenized_datasets = datasets.map(tokenize_function, batched=True)
```

```
# create the train/test datasets and model, and train the model using the Trainer
train_dataset = tokenized_datasets["train"].shuffle(seed=42)
test_dataset = tokenized_datasets["test"].shuffle(seed=42)
```

```
model = AutoModelForSequenceClassification.from_pretrained("bert-base-cased", num_labels=3)
# let's try freezing lower layers
for name, param in model.named_parameters():
    if name.startswith("bert.encoder.layer.0."):
        param.requires_grad = False
```

Some weights of the model checkpoint at bert-base-cased were not used when initializing BertForSequenceClassification:

- This IS expected if you are initializing BertForSequenceClassification from the checkpoint of a model trained on another task (e.g. translation) with a vocabulary size different from the vocabulary of the data (e.g. word embeddings)
- This IS NOT expected if you are initializing BertForSequenceClassification from the checkpoint of a model that you expect to be trained on your task (e.g. adding a new output layer to a pretrained embedding layer)

Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-cased and are newly initialized from random normal distribution. You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

```
training_args = TrainingArguments(output_dir="test_trainer", \
                                  evaluation_strategy='epoch', \
                                  learning_rate=1e-5, \
                                  num_train_epochs=5, \
                                  per_device_train_batch_size=16, \
                                  per_device_eval_batch_size=16 \
                                  #weight_decay=4e-5,\
                                  #max_grad_norm=0.5
                                  )
```

We create a Trainer object with our model, training arguments, training and test_datasets and evaluation function.

```
trainer = Trainer(model=model, \
                  args = training_args, \
```

```

train_dataset=train_dataset, \
eval_dataset = test_dataset, \
compute_metrics = compute_metrics
)

```

Let's compute the class weights

```
from sklearn.utils.class_weight import compute_class_weight
```

```
#compute the class weights
```

```
class_wts = compute_class_weight(class_weight='balanced', classes=np.unique(y_train), y=y_train.to_numpy())
class_wts
```

```
array([ 0.393361 , 14.8125 , 2.56216216])
```

```
from torch import nn
import torch
```

```
class CustomTrainer(Trainer):
```

```

    def compute_loss(self, model, inputs, return_outputs=False):
        labels = inputs.get("labels")
        # forward pass
        outputs = model(**inputs)
        logits = outputs.get("logits")
        # compute custom loss (suppose one has 3 labels with different weights)
        loss_fct = nn.CrossEntropyLoss(weight=torch.tensor([ 0.393361 , 14.8125 , 2.56216216], device=model.device))
        loss = loss_fct(logits.view(-1, self.model.config.num_labels), labels.view(-1))
        return (loss, outputs) if return_outputs else loss

```

```

trainer = CustomTrainer(model=model, \
                        args = training_args, \
                        train_dataset=train_dataset, \
                        eval_dataset = test_dataset, \

```

```
compute_metrics = compute_metrics
)
```

Fine-tune the model by calling train()

```
trainer.train()
```

```
/usr/local/lib/python3.10/dist-packages/transformers/optimization.py:411: FutureWarning: This implementation of AdamW i
warnings.warn(
```

```
[760/760 22:19, Epoch 5/5]
```

Epoch	Training Loss	Validation Loss	Accuracy
1	No log	1.060997	0.846803
2	No log	1.136609	0.179902
3	No log	1.040141	0.843289
4	0.127600	0.941651	0.839775
5	0.127600	0.923681	0.749824

```
TrainOutput(global_step=760, training_loss=0.10901455628244501, metrics={'train_runtime': 1340.8226,
'train_samples_per_second': 9.032, 'train_steps_per_second': 0.567, 'total_flos': 3186303488686080.0, 'train_loss':
0.10901455628244501, 'epoch': 5.0})
```

```
# generate predictions for the test set
predictions = trainer.predict(test_dataset)
```

```
preds = np.array([np.argmax(entry, axis=-1) for entry in predictions[0]])
```

```
f1_metric = evaluate.load('f1')
```

```
results = f1_metric.compute(predictions=preds, references=predictions[1], average='macro')
print(round(results['f1'], 3))
```

0.533

```
f1_metric = evaluate.load('f1')
```

```
results = f1_metric.compute(predictions=preds, references=predictions[1], average='micro')
print(round(results['f1'], 3))
```

0.75

```
print(classification_report(predictions[1], preds))
```

	precision	recall	f1-score	support
0	0.93	0.77	0.85	1205
1	0.50	0.24	0.33	33
2	0.31	0.68	0.43	185
accuracy			0.75	1423
macro avg	0.58	0.57	0.53	1423
weighted avg	0.84	0.75	0.78	1423

```
# derive a confusion matrix
```

```
confusion_matrix(predictions[1], preds)
```

```
array([[933,  5, 267],
       [ 12,  8,  13],
       [ 56,  3, 126]])
```

Since we have an imbalanced data set, we use the weighted average score of 0.78

Generate explanations of the predictions made by the model using XAI techniques that apply to text.

```
!pip install shap
```

Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public/simple/>

Collecting shap

Downloading shap-0.41.0-cp310-cp310-manylinux_2_12_x86_64.manylinux2010_x86_64.whl (572 kB)

572.6/572.6 kB 37.2 MB/s eta 0:00:00

Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (from shap) (1.22.4)

Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-packages (from shap) (1.10.1)

Requirement already satisfied: scikit-learn in /usr/local/lib/python3.10/dist-packages (from shap) (1.2.2)

Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from shap) (1.5.3)

Requirement already satisfied: tqdm>4.25.0 in /usr/local/lib/python3.10/dist-packages (from shap) (4.65.0)

Requirement already satisfied: packaging>20.9 in /usr/local/lib/python3.10/dist-packages (from shap) (23.1)

Collecting slicer==0.0.7 (from shap)

Downloading slicer-0.0.7-py3-none-any.whl (14 kB)

Requirement already satisfied: numba in /usr/local/lib/python3.10/dist-packages (from shap) (0.56.4)

Requirement already satisfied: cloudpickle in /usr/local/lib/python3.10/dist-packages (from shap) (2.2.1)

Requirement already satisfied: llvmlite<0.40,>=0.39.0dev0 in /usr/local/lib/python3.10/dist-packages (from numba->shap)

Requirement already satisfied: setuptools in /usr/local/lib/python3.10/dist-packages (from numba->shap) (67.7.2)

Requirement already satisfied: python-dateutil>=2.8.1 in /usr/local/lib/python3.10/dist-packages (from pandas->shap) (2)

Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas->shap) (2022.7.1)

Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from scikit-learn->shap) (1.2.

Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn->shap)

Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.1->pandas)

Installing collected packages: slicer, shap

Successfully installed shap-0.41.0 slicer-0.0.7

```
import shap
```

```
import transformers
```

Prepare the data set.

```
data = pd.DataFrame({'text':X_train_set['text'],'label':X_train_set['label']})
```

```
data.head(3)
```

	text	label
0	CLAIM, An image of hundreds of people cramped ...	0
1	CLAIM, An image of hundreds of people cramped ...	1

load the model, tokenizer and build a pipeline

```
# load the model and tokenizer
tokenizer = transformers.AutoTokenizer.from_pretrained("bert-base-cased", use_fast=True)
model = transformers.AutoModelForSequenceClassification.from_pretrained("bert-base-cased").cuda()


# build a pipeline object to do predictions
pred = transformers.pipeline("text-classification", model=model, tokenizer=tokenizer, device=0, return_all_scores=True, trunc
```

Some weights of the model checkpoint at bert-base-cased were not used when initializing BertForSequenceClassification:

- This IS expected if you are initializing BertForSequenceClassification from the checkpoint of a model trained on another task (e.g. using `PytorchSeq2LstmTrainer`), since the checkpoint was only for this task and does not contain parameters for this task.
- This IS NOT expected if you are initializing BertForSequenceClassification from the checkpoint of a model that you expect (e.g. using `BertForSequenceClassification`), since the model was trained on this task (or a similar task) and should have initialized all weights.

Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-cased and are newly initialized from random normal distribution. You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

`return_all_scores` is now deprecated, if want a similar functionality use `top_k=None` instead of `return_all_scores=True`. Xformers is not installed correctly. If you want to use memory_efficient_attention to accelerate training use the following pip install xformers.



Create an explainer object for the pipeline.

```
explainer = shap.Explainer(pred)
```

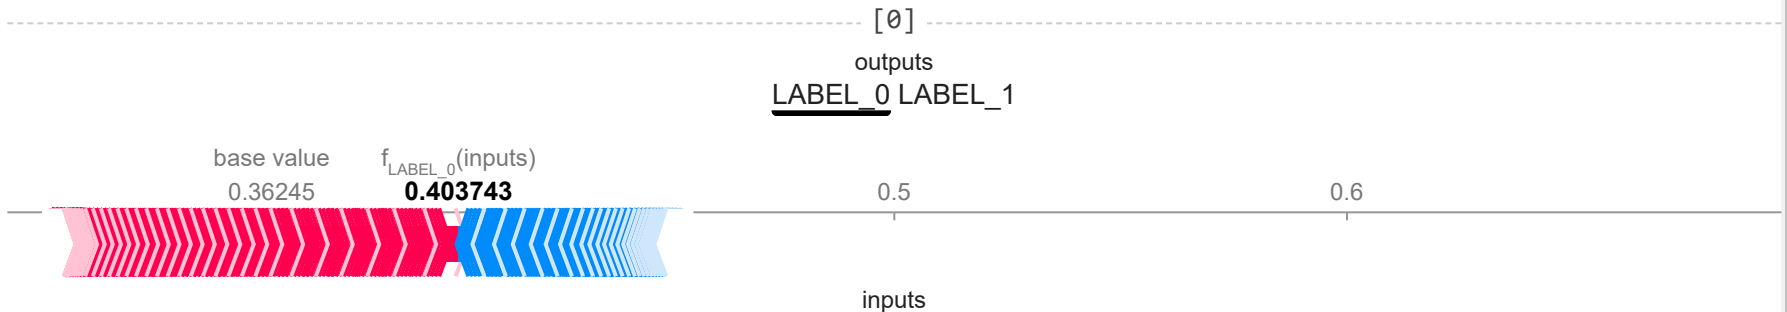
Compute the Shap values

```
shap_values = explainer(data['text'][:3])
```


Token indices sequence length is longer than the specified maximum sequence length for this model (696 > 512). Running
You seem to be using the pipelines sequentially on GPU. In order to maximize efficiency please use a dataset

Visualizing the impact on the output classes:

```
shap.plots.text(shap_values)
```

CLAIM, An image of hundreds of people cramped in an alley is being circulated to claim that it shows lockdown violation in Gujarat. The image is being used to insinuate that these are migrant workers and homeless people who have nowhere to go., The claim reads: "ये फोटो #गुजरात का बताया जा रहा है। इनको भुल जाओ और जमाती में फंसे रहे? इसको कहते हैं राजनीति असल मुद्दा से नजर हटाना और धर्म और जात पात में उलझा देना हाय रे हमारे भारत वासियों |" Translation: This is from Gujarat. Should we ignore and only talk about the Jamaatis? This is called the politics of diverting from real issues and fooling people in the name of caste and religion., As on 16 April, Gujarat has 695 coronavirus cases with 30 recorded deaths., (Click here for live updates on COVID-19. Also visit Quint Fit for comprehensive coverage on the impact of the coronavirus pandemic.), The image is also being circulated on Twitter with the same claim., TRUE OR FALSE?, The image is being circulated with a misleading claim. It does not show the plight of migrant workers or homeless people in Gujarat but shows Bangladeshi refugees in Malaysia waiting to go back home., WHAT WE FOUND, We conducted a Google reverse image search and found a report dated 25 December 2019 by a Bangladesh based news website, Daily Fulki which carried the same image., "As per the report, thousands of immigrants crowded at Malaysia's immigration office after the country announced to facilitate the Back for Good program for repatriation of illegal foreigners.", Another report by JagoNews24 corroborated the details. "The migrants from different countries have been waiting a long time before opening the counter at 9 am, early or before midnight. Thousands of immigrants are waiting in line, but on an average, 8 special passes are being issued every day from an immigration office. The rest have to be in line in anticipation of the next day," the report stated., Hence, an old image of Bangladeshi immigrants in Malaysia is being shared to claim that it shows a violation of COVID-19 lockdown in Gujarat., You can read all our fact-checked stories here., "(Not convinced of a post or information you came across online and want it verified? Send us the details on WhatsApp at 9643651818, or e-mail it to us at webqoof@thequint.com and well fact-check it for you. You can also read all our fact-checked stories here.)", "Well get through this! Meanwhile, heres all you need to know about the Coronavirus outbreak to keep yourself safe, informed, and updated.", The Quint is now available on Telegram & WhatsApp too, Click here to join., Stay tuned with our weekly recap of what's hot & cool by The Quint.



[more info, as more people are infected with covid - 19, the disease caused by the novel coronavirus, many are wondering : am i next?, the


odds may not be in our favor, according to a recent facebook post., " harvard scientists just predicted that it may take many months to find cure for coronavirus, " the march 10 post says. " according to them, virus is spreading so fast that it will infect 70 % of humanity this year. ", this post was flagged as part of facebook ' s efforts to combat false news and misinformation on its news feed. (read more about our partnership with facebook.), spokespeople for harvard university did not immediately reply to politifact ' s emails asking about the post., but news stories lend some credence to the facebook post ' s claims — and provide more information that the post fails to mention., on march 11, the world health organization declared covid - 19 a pandemic that has infected nearly 120, 000 people worldwide. more than 4, 000 people have died and eight countries, including the united states, are currently each reporting more than 1, 000 cases., mark lipsitch, an epidemiologist at harvard ' s center for communicable disease dynamics and an expert on viruses, told cbs news, among other news outlets, that 40 % to 70 % of the world ' s adult population could be infected in the coming year., there are an estimated 7. 6 billion people in the world, and an estimated 329 million in the united states, according to the u. s. census bureau. adults make up about three - quarters of the u. s. population., " that is a projection, " lipsitch said during the march 2 interview, " so we will find out if that ' s accurate as things go on. it is the best estimate that i ' ve been able to make based on a combination of the mathematical models that we use to track and predict epidemics. ", other estimates are similar. the

Visualizing the impact on a single class.

researchers need more data to get a better idea of how far the coronavirus could spread. " the models will be refined as more information

```
# visualizing the impact of class 1
shap.plots.text(shap_values[:, :, 1])
```


Creating an ndarray from ragged nested sequences (which is a list-or-tuple of lists-or-tuples-or ndarrays with different lengths)




CLAIM, An image of hundreds of people cramped in an alley is being circulated to claim that it shows lockdown violation in Gujarat. The image is being used to insinuate that these are migrant workers and homeless people who have nowhere to go. The claim reads: "ये फोटो #गुजरात का बताया जा रहा है। इनको भुल जाओ और जमाती में फंसे रहे? इसको कहते हैं राजनीति असल मुद्दा से नजर हटाना और धर्म और जात पात में उलझा देना हाथों हमारे भारत वासियों।" Translation: This is from Gujarat. Should we ignore and only talk about the Jamaatis? This is called the politics of diverting from real issues and fooling people in the name of caste and religion. As on 16 April, Gujarat has 695 coronavirus cases with 30 recorded deaths. (Click here for live updates on COVID-19. Also visit Quint Fit for comprehensive coverage on the impact of the coronavirus pandemic.) The image is also being circulated on Twitter with the same claim. TRUE OR FALSE? The image is being circulated with a misleading claim. It does not show the plight of migrant workers or homeless people in Gujarat but shows Bangladeshi refugees in Malaysia waiting to go back home. WHAT WE FOUND, We conducted a Google reverse image search and found a report dated 25 December 2019 by a Bangladesh based news website, Daily Fulki which carried the same image. "As per the report, thousands of immigrants crowded at Malaysia's immigration office after the country announced to facilitate the Back for Good program for repatriation of illegal foreigners." Another report by JagoNews24 corroborated the details. "The migrants from different countries have been waiting a long time before opening the counter at 9 am, early or before midnight. Thousands of immigrants are waiting in line, but on an average, 8 special passes are being issued every day from an immigration office. The rest have to be in line in anticipation of the next day," the report stated. Hence, an old image of Bangladeshi immigrants in Malaysia is being shared to claim that it shows a violation of COVID-19 lockdown in Gujarat. You can read all our fact-checked stories here. (Not convinced of a post or information you came across online and want it verified? Send us the details on WhatsApp at 9643651818, or e-mail it to us at webqoof@thequint.com and we'll fact-check it for you. You can also read all our fact-checked stories here.) "Well get through this! Meanwhile, here's all you need to know about the Coronavirus outbreak to keep yourself safe, informed, and updated." The Quint is now available on Telegram & WhatsApp too, Click here to join. Stay tuned with our weekly recap of what's hot & cool by The Quint.



```
# visualizing the impact of class 0
shap.plots.text(shap_values[:, :, 0])
```



CLAIM, An image of hundreds of people cramped in an alley is being circulated to claim that it shows lockdown violation in Gujarat. The image is being used to insinuate that these are migrant workers and homeless people who have nowhere to go., The claim reads: "ये फोटो #गुजरात का बताया जा रहा है। इनको भुल जाओ और जमाती में फंसे रहे? इसको कहते हैं राजनीति असल मुद्दा से नजर हटाना और धर्म और जात पात में उलझा देना हाय रे हमारे भारतवासियों।" Translation: This is from Gujarat. Should we ignore and only talk about the Jamaatis? This is called the politics of diverting from real issues and fooling people in the name of caste and religion., As on 16 April, Gujarat has 695 coronavirus cases with 30 recorded deaths., (Click here for live updates on COVID-19. Also visit Quint Fit for comprehensive coverage on the impact of the coronavirus pandemic.), The image is also being circulated on Twitter with the same claim., TRUE OR FALSE?, The image is being circulated with a misleading claim. It does not show the plight of migrant workers or homeless people in Gujarat but shows Bangladeshi refugees in Malaysia waiting to go back home., WHAT WE FOUND, We conducted a Google reverse image search and found a report dated 25 December 2019 by a Bangladesh based news website, Daily Fulki which carried the same image., "As per the report, thousands of immigrants crowded at Malaysia's immigration office after the country announced to facilitate the Back for Good program for repatriation of illegal foreigners.", Another report by JagoNews24 corroborated the details. "The migrants from different countries have been waiting a long time before opening the counter at 9 am, early or before midnight. Thousands of immigrants are waiting in line, but on an average, 8 special passes are being issued every day from an immigration office. The rest have to be in line in anticipation of the next day," the report stated., Hence, an old image of Bangladeshi immigrants in Malaysia is being shared to claim that it shows a violation of COVID-19 lockdown in Gujarat., You can read all our fact-checked stories here., "(Not convinced of a post or information you came across online and want it verified? Send us the details on WhatsApp at 9643651818, or e-mail it to us at webqoof@thequint.com and we'll fact-check it for you. You can also read all our fact-checked stories here.)", "Well get through this! Meanwhile, here's all you need to know about the Coronavirus outbreak to keep yourself safe, informed, and updated.", The Quint is now available on Telegram & WhatsApp too, Click here to join., Stay tuned with our weekly recap of what's hot & cool by The Quint.



[more info, as more people are infected with covid - 19, the disease caused by the novel coronavirus, many are wondering : am i next?, the odds may not be in our favor, according to a recent facebook post., " harvard scientists just predicted that it may take many months to find cure for coronavirus, " the march 10 post says. " according to them, virus is spreading so fast that it will infect 70 % of humanity this year. ", this post was flagged as part of facebook ' s efforts to combat false news and misinformation on its news feed. (read more about our partnership with facebook.), spokespeople for harvard university did not immediately reply to politifact ' s emails asking about the post., but news stories lend some credence to the facebook post ' s claims — and provide more information that the post fails to mention., on march 11, the world health organization declared covid - 19 a pandemic that has infected nearly 120, 000 people worldwide. more than 4, 000 people have died and eight

countries, including the united states, are currently each reporting more than 1, 000 cases., mark lipsitch, an epidemiologist at harvard ' s center for communicable disease dynamics and an expert on viruses, told cbs news, among other news outlets, that 40 % to 70 % of the world ' s adult population could be infected in the coming year., there are an estimated 7. 6 billion people in the world, and an estimated 329 million in the united states, according to the u. s. census bureau. adults make up about three - quarters of the u. s. population., " that is a projection, " lipsitch said during the march 2 interview, " so we will find out if that ' s accurate as things go on. it is the best estimate that i ' ve been able to make based on a combination of the mathematical models that we use to track and predict epidemics. ", other estimates are similar. the guardian reported on feb. 11 that 60 % of the global population could be infected if the virus is unchecked. on feb. 13, bloomberg reported that ira longini, an adviser to the world health food organization, said two - thirds of the world ' s population could catch covid - 19., but the publication also quoted david heymann, an infectious disease expert at the london school of hygiene & tropical medicine, explaining that researchers need more data to get a better idea of how far the coronavirus could spread., " the models will be refined as more information comes to light. " he said.. our ruling. the facebook post says harvard scientists say the coronavirus is " spreading so fast that it will infect 70 %

Plotting the words that impact class 0

coronavirus may infect up to 70 % of world ' s population, expert warns, march 2, 2020, the hill, harvard scientist : coronavirus pandemic likely

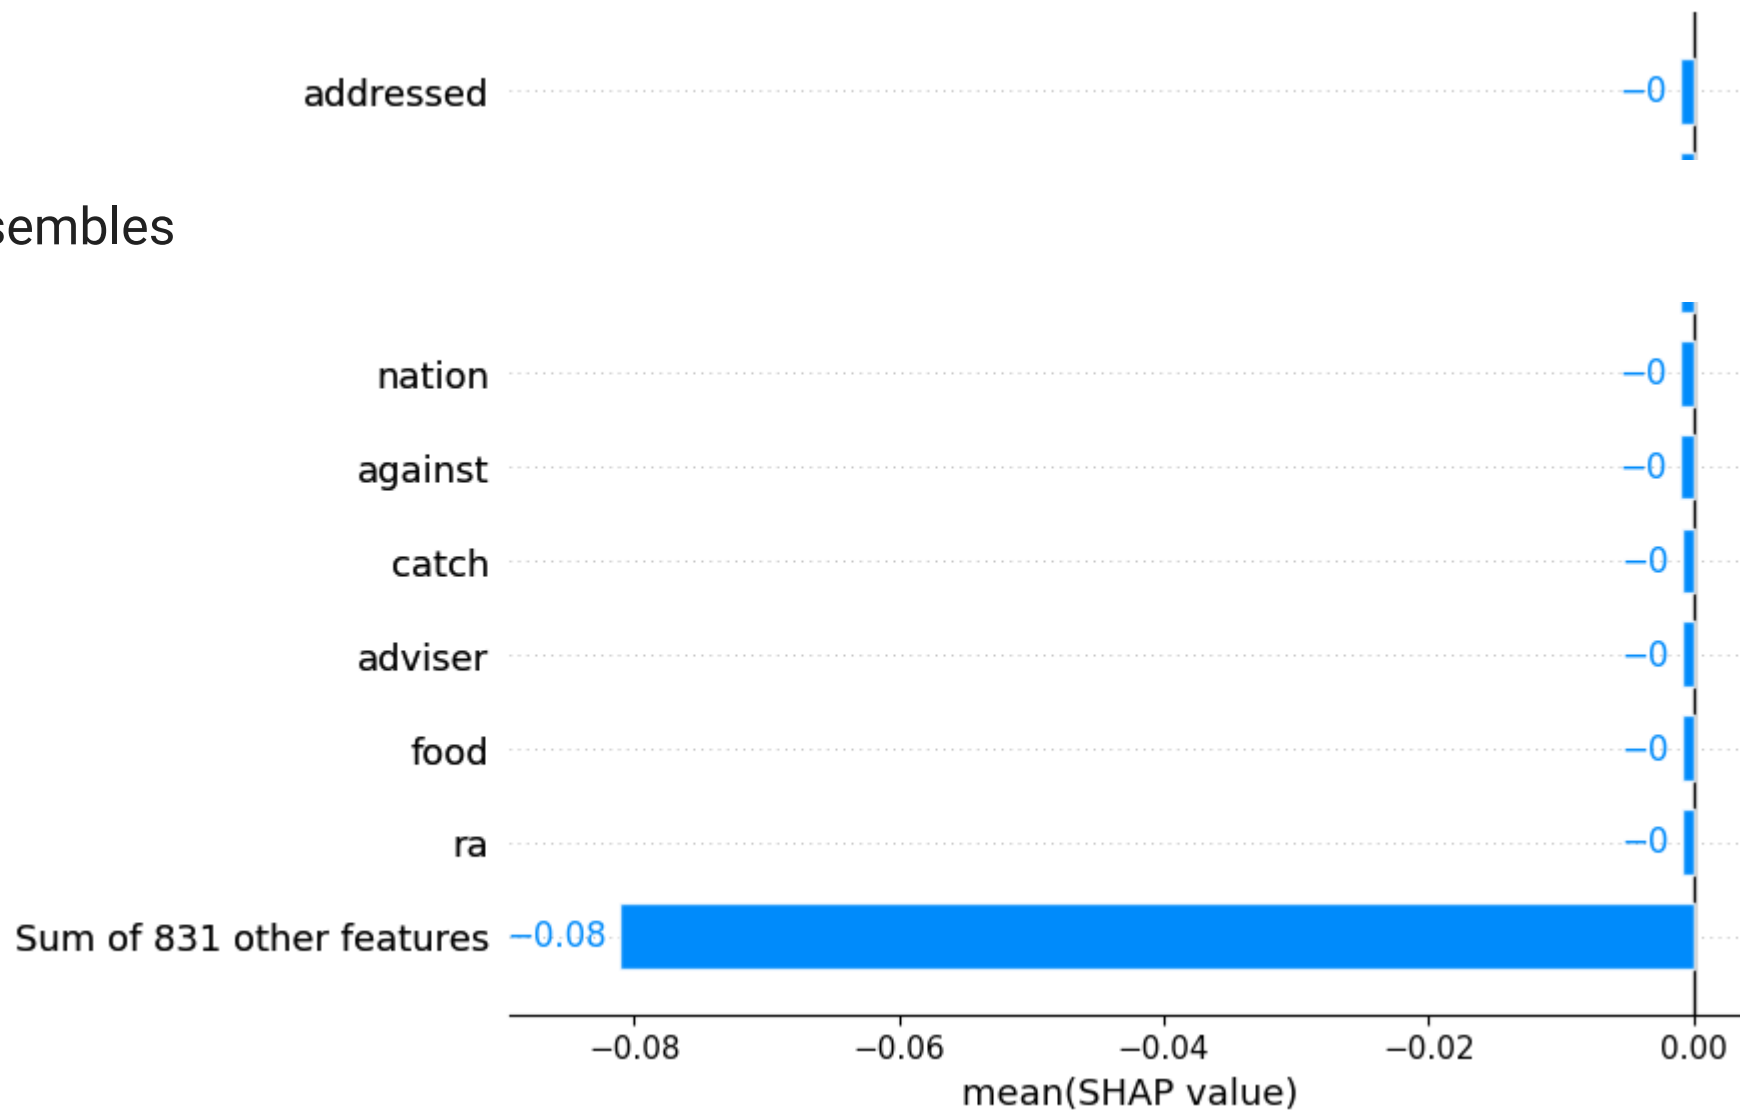
```
shap.plots.bar(shap_values[:, :, 0].mean(0))
```



Plotting the words that impact class 1

```
.. ■ .  
shap.plots.bar(shap_values[:, :, 1].mean(0))
```

Ensembles



[Colab paid products](#) - [Cancel contracts here](#)

